

Math 338 Lab 24 Fall 2019

In this lab, we will investigate the importance of evaluating whether the assumptions of the Simple Linear Regression Model are met. We will work with a dataset called *anscombe*, which actually contains four (x, y) datasets that you will find have the same statistical properties, but are VERY different.

```
as <- anscombe
library(ggplot2)
old_x1y1_plot <- ggplot(as, mapping = aes(x = x1, y = y1)) + geom_point()
lm_x1y1 <- lm(y1 ~ x1, data = as)
coef_x1y1 <- coef(lm_x1y1)
new_x1y1_plot <- old_x1y1_plot + geom_abline(intercept = coef_x1y1[1], slope = coef_x1y1[2])
print(new_x1y1_plot)
```

Question 1 Insert the final plot below.




Figure 1: Graph One




Figure 2: Graph Two




Figure 3: Graph Three




Figure 4: Graph Four

Question 2 Use the *summary()* command to learn more about your linear model. Insert the *Coefficients:* table from the summary below.

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.0001	1.1247	2.667	0.02573 *
x1	0.5001	0.1179	4.241	0.00217 **

Question 3: Write out the equation of the least-squares regression line for x1 and y1 using the intercept and slope estimates given in the output. Round your estimates to 2 decimal places and remember to use the actual names of the variables when writing out the equation.

$$y_1 = 3.00 + 0.50x_1$$

Question 4 What is the observed value of the t-statistic for the hypothesis test for the slope parameter? Show how you can compute it from the values of b_1 and SE_{b_1} .

- The observed value of the t-statistic is estimate error over standard error which is going to be $\frac{0.5001}{0.1179} = 4.241$

Question 5 What is the p-value associated with the hypothesis test for the slope? At the 5% significance level, what can you conclude about the slope?

- The p-value is going to be 0.00217 and at an α value of 0.05, we can conclude that the slope is not zero.

Question 6 What is the value of the coefficient of determination?

- Coefficient of determination is 0.6665

Now, let's look at the residual plot.

Question 7 Insert the plot below. Do the residuals look healthy (roughly scattered around 0 with no obvious pattern or “fanning”)?

- The residual is healthy because it does not rely on \hat{y}




Figure 5: Graph One

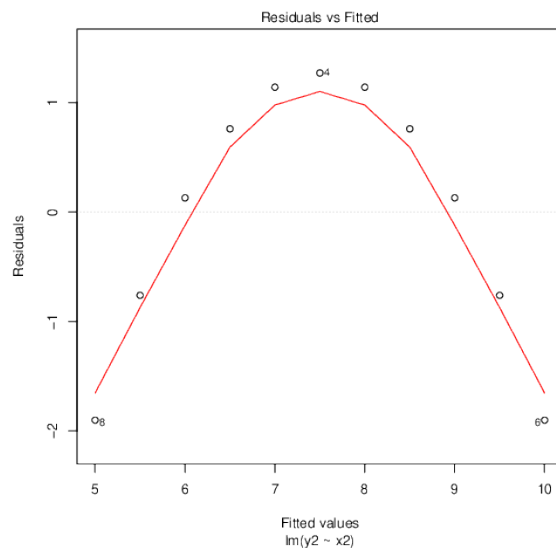




Figure 6: Graph Two



We can also look at the normal q-q plot to assess normality.

Question 8 Insert the plot below. Do the residuals look approximately normal?

- The residual looks approximately normal since it is in fact linear.




Figure 9: Graph One




Figure 10: Graph Two




Figure 11: Graph Three




Figure 12: Graph Four

Question 9 Repeat questions 1-8 for the other three datasets: x_2 - y_2 , x_3 - y_3 , and x_4 - y_4 . (You can put your answers with the corresponding question.) You should find that using linear regression with x_2 - y_2 is obviously a bad idea. **Should linear regression be used to model the x_3 - y_3 relationship? What about x_4 - y_4 ? Justify your answer.**

- For x_3 - y_3 , yes the regression line should be used since the points are close to the line even though there are outliers
- For x_4 - y_4 , however, the regression line should not be used because there is a vertical line, implying there is no slope.

Question 10 Why is it important to look at the scatterplot and residual plots? Why can't we just look at the least-squares regression line, r^2 , and the p-value?

- By looking at the data plotted out, you can see if there is a need to use a linear regression line or not. Both the R^2 and the p-value can be the same for different datasets and is not consistent.