

Math 338 Lab Assignment 25 Fall 2019

For this lab, we will perform several different types of inference for simple linear regression. We will work with two variables from the *longley* dataset in RStudio, which describe a total of seven macroeconomic variables observed yearly from 1947 to 1962.

We could explore many different relationships among the seven variables. For now, we will focus on the Gross National Product (GNP) as a linear function of the number of unemployed (Unemployed).

```
summary(longley)
```

Question 1 Create a scatterplot of the two variables (with the response on the y-axis and the predictor on the x-axis) using the *ggplot()* and *geom_point()* commands. Insert the plot below.




Figure 1: Scatter Plot

Question 2 Does there appear to be a linear relationship between GNP and Unemployed? If so, comment on the strength of the relationship.

Now, let's obtain the linear model for these data.

There appears to be a linear relationship because the data points seem to be equally spaced between the regression line. However it is not particularly strong.

```
longley_lm <- lm(GNP ~ Unemployed, data = data)
summary(longley_lm)
```

Question 3 Copy and paste the **Coefficients:** table below. Use the estimates in the table to write the least-squares regression equation.

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	182.4558	75.1778	2.427	0.0293	*
Unemployed	0.6427	0.2265	2.838	0.0132	*

$y = 0.6427x + 182.4558$

Question 4 Using your least-squares regression equation, predict the GNP when Unemployment is at 187.

$$x = 187$$

$$y = 0.6427 \times 187 + 182.4588 = 302.6407$$

Question 5 Write the null and alternative hypothesis for the (NHST-framework) hypothesis test for the slope parameter. Also, write the (null) hypothesis associated with the ANOVA table.

NHST

- $H_0: \beta_1 = 0$ which means there is no linear relationship
- $H_a: \beta_1 \neq 0$ which means there is a linear relationship

ANOVA

- $\mu_{y|x} = \beta_0 + \beta_1 x$
- $\therefore \beta_1 = 0$

Question 6 What type of test statistic can be read from the **Coefficients:** table, and what is its observed value in the test for slope? What is the p-value? At the 5% significance level, what can you conclude about the slope?

- We are reading a t-statistic and its observed value is 2.838 and the p-value is 0.0132. Since our p-value is less than α , we can reject the null hypothesis

Now let's get the ANOVA table for this linear model.

```
anova(longley_lm)
```

Question 7 Copy and paste the table below. Note that the summary only includes the first two rows of the ANOVA table; it does not include the bottom row (Total).

Response: GNP

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Unemployed	1	54109	54109	8.0518	0.01317 *
Residuals	14	94081	6720		

Question 8 What type of test statistic can be read from the ANOVA table, what distribution does it come from (include the degrees of freedom), and what is its value in this ANOVA test? What is the p-value, according to the ANOVA table? At the 5% significance level, what can you conclude about the appropriateness of the linear model?

- This would be a F-Statistic that comes from a F-Distribution with $F \sim (1, 14)$. With a p-value of 0.01317, we can reject the null hypothesis because it is less than our significance level $\alpha = 0.05$

Now we're going to make predictions using our dataset. Recall that we can predict by computing \hat{y} , by computing a confidence interval for μ_y , or by computing a prediction interval for y .

First, we have to set up a new data frame containing the predictor values we want to make predictions at.

```
new_data_frame <- data.frame(Unemployed = c(187, 200, 308))
```

To make point estimate \hat{y} predictions, we use the **predict** command. If we then want to include a confidence interval or prediction interval centered at \hat{y} , include the **interval** argument to tell it which type of interval to create.

```
longley_predictions <- predict(longley_lm, newdata = new_data_frame)
```

```
longley_predictions_PI <- predict(longley_lm, newdata = new_data_frame, interval = "confidence", level = 0.95)
```

```
longley_predictions_PI <- predict(longley_lm, newdata = new_data_frame, interval = "confidence", level = 0.95)
```

To make things easier to read, we then combine our predictor values and prediction values into a single data frame.

```

longley_CI_df <- data.frame(new_data_frame, longley_predictions_CI)
longley_PI_df <- data.frame(new_data_frame, longley_predictions_PI)

print(longley_CI_df)
print(longley_PI_df)

```

Question 9 In the output, find the 95% CI for the mean GNP when Unemployment is at 187 and a 95% PI for the actual GNP when Unemployment is at 187. Label which interval is which.

- Confidence Interval: (224.7679, 380.5234)
- Prediction Interval: (110.3486, 494.9427)

Create the residual plot and normal q-q plot:

```
plot(longley_lm, which = c(1, 2))
```

Question 10 Insert the residual plot and normal q-q plot below. Are any assumptions necessary for inference (residuals are independent and identically distributed normally with mean 0 and constant variance) clearly violated? If so, explain which ones.




Figure 2: Residual Plot

- The time Δ in this data set is one year, which means the residuals in each year will affect each other because not enough time has passed.